

Stochastic Degradation Aware Optimal Fleet management with Application to Electric Vehicles

A Comprehensive Reference Document

L. Fiaschi
ETH Zürich
lfiaschi@ethz.ch

May 2026

Abstract

This is a project scoping where I collect all the needed information to write an article. The goal is to develop a general purpose model and its algorithmic solver to address the management of a fleet during a certain horizon. Each member of the fleet is described as a collection of components. Each day, a member can be sent to serve a route, to maintenance (partial or full repair of each component), or left idle. The service of a route damages the components of the fleet member. Each component has its own stochastic degradation model, and the actual degradation changes according to the route served. We show a concrete application to the management of electric-vehicle fleets, such as buses and cars.

Contents

1	Overview of the problem	3
1.1	Scheduling and maintenance variables	3
1.2	Damage dynamics	3
1.3	Stochastic loop constraint	5
1.4	Cost function	6
1.5	Abstract problem formulation	6
2	Literature review	7
2.1	Stochastic degradation and fatigue models	7
2.1.1	Damage-accumulation processes	7
2.1.2	Discrete-time formulation	8
2.1.3	Remaining-life models: the S–N curve and Palmgren–Miner rule	8
2.1.4	Conservative surrogates for the empirical damage-increment distribution	10
2.1.5	The loop constraint under the Route-I bounds	14
2.2	Stochastic orderings	15
2.2.1	Stochastic ordering within each degradation family	16
2.2.2	Loop constraint for each degradation model	17
2.3	Maintenance strategies	19
2.3.1	Taxonomy of repair actions	19
2.3.2	Virtual age models	19
2.3.3	Reduction-based models	20
2.3.4	Connection to the fleet management formulation	20
3	Focus on electric vehicles	21
3.1	Degradation models for the shortlisted components	24

4	Optimization problem for a fleet of electric vehicles	27
4.1	Cost function	28
5	Implementation details	28

1 Overview of the problem

We consider a fleet of F members, each composed of L components subject to stochastic degradation. The planning horizon spans $2H$ discrete time steps (days). At each step, every fleet member is assigned to exactly one activity: serving one of M available missions, undergoing maintenance, or remaining idle. Serving a mission damages the components of the assigned member according to a stochastic degradation process that depends on both the mission and the component. On a maintenance day, each component may independently undergo imperfect repair (which partially reduces accumulated damage), full replacement (which resets the component to a like-new state), or no intervention at all.

The degradation state of component ℓ of fleet member i at time k is described by a random variable $D_{i\ell k}$, whose distribution belongs to a parametric family specified per component. A component is considered to have failed when $D_{i\ell k}$ exceeds a threshold τ ; the schedule must guarantee that the probability of this event remains below ε at every time step. To ensure that the schedule can be repeated indefinitely, a *loop constraint* requires that the degradation state at the end of the horizon stochastically dominates (in the sense of a lower probability of failure) the state at the mid-point: the second half of the schedule must leave the fleet in at least as good a condition as the first half did.

Table 1 collects the notation used throughout this document. All variables are general and independent of any specific degradation model.

1.1 Scheduling and maintenance variables

The schedule is encoded by three families of binary decision variables. The *assignment variable* x_{ijk} determines which activity fleet member i performs on day k : either serving mission $j \geq 1$ or spending the day in the depot ($j = 0$), either maintenance or idle. Each member performs at most one activity per day,

$$\sum_{j=0}^M x_{ijk} = 1 \quad \forall i, k, \quad (1)$$

and each mission must be covered by exactly one member,

$$\sum_{i=1}^F x_{ijk} = 1 \quad \forall j \geq 1, k. \quad (2)$$

Since $F > M$, at least $F - M$ members are available for maintenance or idling at every step.

On a maintenance day ($x_{i0k} = 1$), the intervention on each component is controlled independently by two additional binary variables. The *repair variable* $x_{i\ell k}^m$ indicates whether component ℓ undergoes imperfect repair, and the *replacement variable* $x_{i\ell k}^r$ indicates full replacement. Both require a maintenance day and are mutually exclusive:

$$x_{i\ell k}^m \leq x_{i0k}, \quad x_{i\ell k}^r \leq x_{i0k}, \quad x_{i\ell k}^m + x_{i\ell k}^r \leq 1 \quad \forall i, \ell, k. \quad (3)$$

A component on a maintenance day with $x_{i\ell k}^m = x_{i\ell k}^r = 0$ receives no intervention (stays idle); its degradation state carries over unchanged.

1.2 Damage dynamics

The degradation of each component evolves in discrete time according to the activity assigned to the fleet member that carries it. We write $\Delta D_{i\ell k}^{(j)}$ for the *random damage increment* sustained by component ℓ of member i when mission j is served on day k . The distribution of each increment is assumed known (from the component's degradation model) and may depend on i, ℓ, j , and k . On days when the fleet member is idle or undergoing maintenance, no damage is accumulated.

Table 1: Notation and variable definitions.

Symbol	Description
<i>Index sets</i>	
$i \in \{1, \dots, F\}$	Fleet member
$j \in \{0, 1, \dots, M\}$	Activity ($j = 0$: maintenance day; $j \geq 1$: mission)
$k \in \{1, \dots, 2H\}$	Time step (day)
$\ell \in \{1, \dots, L\}$	Component
<i>Scalar and fleet-level parameters</i>	
F	Number of fleet members
H	Half-horizon length (the full planning window spans $2H$ days)
M	Number of missions (must satisfy $F > M$)
L	Number of components per fleet member
τ	Degradation failure threshold (component-level)
ε	Reliability threshold: $\Pr(D_{i\ell k} > \tau) \leq \varepsilon$ for all i, ℓ, k
<i>Cost functions</i>	
C_M	Maintenance cost: $\{x_{i0k}\}_{i,k} \rightarrow \mathbb{R}_{\geq 0}$
C_R	Repair cost: $\{z_{i\ell k}\}_{i,\ell,k} \rightarrow \mathbb{R}_{\geq 0}$
C_{rep}	Replacement cost: $\{x_{i\ell k}^r\}_{i,\ell,k} \rightarrow \mathbb{R}_{\geq 0}$
C_D	Damage regularisation cost: $u \rightarrow \mathbb{R}_{\geq 0}$
<i>Component-level parameters</i>	
$\rho_{i\ell}$	Repair efficiency $\in (0, 1]$
<i>Initial and reference values</i>	
$D_{i\ell}^{\text{start}}$	Initial degradation state of component ℓ of member i (known distribution; must satisfy $\Pr(D_{i\ell}^{\text{start}} > \tau) \leq \varepsilon$)
$D_{i\ell}^{\text{new}}$	Post-replacement degradation state (known distribution; typically near-zero damage)
<i>Stochastic state</i>	
$D_{i\ell k}$	Random variable: degradation of component ℓ of member i at time k
$\Delta D_{i\ell k}^{(j)}$	Random damage increment to component ℓ of member i when mission j is served on day k
<i>Decision variables</i>	
$x_{ijk} \in \{0, 1\}$	1 if member i is assigned to activity j on day k
$x_{i\ell k}^m \in \{0, 1\}$	1 if component ℓ of member i undergoes imperfect repair on day k
$x_{i\ell k}^r \in \{0, 1\}$	1 if component ℓ of member i is replaced on day k
$z_{i\ell k} \geq 0$	Degradation removed by imperfect repair (in expected-damage units)
$u \geq 0$	Damage regularisation variable

At each time step, exactly one of the following mutually exclusive cases applies to every component (i, ℓ) :

- (i) **Mission.** Member i serves mission $j \geq 1$ ($x_{ijk} = 1$). The degradation grows by a mission- and component-dependent random increment:

$$D_{ilk} = D_{il,k-1} + \Delta D_{ilk}^{(j)}. \quad (4)$$

- (ii) **Imperfect repair.** The member is on a maintenance day ($x_{i0k} = 1$) and component ℓ undergoes imperfect repair ($x_{ilk}^m = 1$). The repair reduces the accumulated damage by a fraction determined by the repair efficiency ρ_{il} and the specific maintenance policy in use. Abstractly:

$$D_{ilk} = \mathcal{R}(D_{il,k-1}, \rho_{il}), \quad (5)$$

where $\mathcal{R}$ denotes the repair operator. See Section 2.3 for more details. The degradation removed by the repair action defines the repair cost variable:

$$z_{ilk} = \mathbb{E}[D_{il,k-1}] - \mathbb{E}[D_{ilk}], \quad \text{when } x_{ilk}^m = 1. \quad (6)$$

- (iii) **Full replacement.** The member is on a maintenance day and component ℓ is replaced ($x_{ilk}^r = 1$). The degradation is reset:

$$D_{ilk} = D_{il}^{\text{new}}. \quad (7)$$

- (iv) **Idle or maintenance day with no intervention on this component.** The degradation is unchanged:

$$D_{ilk} = D_{il,k-1}. \quad (8)$$

Cases (ii) and (iii) are mutually exclusive: a component cannot be both repaired and replaced on the same day. Both require that the member be assigned to a maintenance day ($x_{i0k} = 1$). Since exactly one case applies at each step, the dynamics can be written as a single equation:

$$D_{ilk} = D_{il,k-1}(1 - x_{ilk}^m)(1 - x_{ilk}^r) + \sum_{j=1}^M x_{ijk} \Delta D_{ilk}^{(j)} + \mathcal{R}(D_{il,k-1}, \rho_{il}) x_{ilk}^m + D_{il}^{\text{new}} x_{ilk}^r, \quad (9)$$

with initial condition $D_{il,0} = D_{il}^{\text{start}}$.

1.3 Stochastic loop constraint

The planning horizon is split into two halves of length H . For the schedule to be repeatable indefinitely, the fleet must emerge from the second half in no worse condition than it had at the midpoint. We formalise this by requiring that the probability of failure at the end of the horizon does not exceed the probability of failure at the midpoint, component by component:

$$\Pr(D_{il,2H} > \tau) \leq \Pr(D_{il,H} > \tau) \quad \forall i, \ell. \quad (10)$$

Together with the reliability constraint $\Pr(D_{ilk} > \tau) \leq \varepsilon$ enforced at every step, condition (10) ensures that the second half of the schedule does not consume more reliability margin than the first half created. Concatenating copies of the schedule therefore yields an infinite-horizon policy whose failure probability never exceeds ε .

1.4 Cost function

The objective is to minimise the total cost of operating the fleet over the planning horizon. Four terms contribute to the cost:

$$\min \quad C_M(\{x_{i0k}\}_{i,k}) + C_R(\{z_{ilk}\}_{i,\ell,k}) + C_{\text{rep}}(\{x_{ilk}^r\}_{i,\ell,k}) + C_D(u). \quad (11)$$

All four cost functions are assumed to be non-negative and non-decreasing in each of their arguments: scheduling more maintenance days, removing more degradation, replacing more components, or tolerating higher aggregate damage never decreases the cost.

The first term charges for every maintenance day scheduled across the fleet; it penalises excessive depot usage. The second term accounts for imperfect repair: whenever a component is repaired, the cost grows with the amount of degradation removed z_{ilk} . The third term charges for every full replacement.

The last term is a damage regularisation penalty. The variable u captures the worst-case aggregate expected damage across all fleet members and time steps:

$$u = \max_{i,k} \phi(\mathbb{E}[D_{i1k}], \dots, \mathbb{E}[D_{iLk}]), \quad (12)$$

where $\phi: \mathbb{R}^L \rightarrow \mathbb{R}_{\geq 0}$ is a non-decreasing aggregation function (e.g. the sum or the sum of squares of the expected component damages). The cost controls the trade-off between minimising operational costs and spreading damage evenly across the fleet.

1.5 Abstract problem formulation

Collecting the objective (11), the scheduling constraints (1)–(3), the degradation dynamics (9), the reliability constraint, and the loop constraint (10), the fleet management problem reads:

$$\min_{x, x^m, x^r, z, u} \quad C_M(\{x_{i0k}\}) + C_R(\{z_{ilk}\}) + C_{\text{rep}}(\{x_{ilk}^r\}) + C_D(u) \quad (13a)$$

$$\text{s.t.} \quad \sum_{j=0}^M x_{ijk} = 1 \quad \forall i, k, \quad (13b)$$

$$\sum_{i=1}^F x_{ijk} = 1 \quad \forall j \geq 1, k, \quad (13c)$$

$$x_{ilk}^m + x_{ilk}^r \leq x_{i0k} \quad \forall i, \ell, k, \quad (13d)$$

$$D_{ilk} = D_{i\ell,k-1}(1 - x_{ilk}^m)(1 - x_{ilk}^r) + \sum_{j=1}^M x_{ijk} \Delta D_{ilk}^{(j)} + \mathcal{R}(D_{i\ell,k-1}, \rho_{i\ell}) x_{ilk}^m + D_{i\ell}^{\text{new}} x_{ilk}^r \quad \forall i, \ell, k, \quad (13e)$$

$$\Pr(D_{ilk} > \tau) \leq \varepsilon \quad \forall i, \ell, k, \quad (13f)$$

$$\Pr(D_{i\ell,2H} > \tau) \leq \Pr(D_{i\ell,H} > \tau) \quad \forall i, \ell, \quad (13g)$$

$$u = \max_{i,k} \phi(\mathbb{E}[D_{i1k}], \dots, \mathbb{E}[D_{iLk}]), \quad (13h)$$

$$x_{ijk}, x_{ilk}^m, x_{ilk}^r \in \{0, 1\}, \quad z_{ilk} \geq 0, \quad u \geq 0. \quad (13i)$$

Problem (13) is a stochastic combinatorial optimisation problem. It makes no commitment to a specific degradation model (the distribution of $\Delta D_{ilk}^{(j)}$), repair policy (the operator $\mathcal{R}$), aggregation function (ϕ), or cost structure (the functions $C_M, C_R, C_{\text{rep}}, C_D$). All of these are instantiated when the formulation is specialised to a concrete fleet (Section 4).

2 Literature review

2.1 Stochastic degradation and fatigue models

Stochastic models of component degradation fall into two broad families. *Damage-accumulation models* track a cumulative damage variable $X(t)$ whose distribution evolves over time; failure occurs when $X(t)$ exceeds a threshold. *Remaining-life models* characterise the number of load cycles a component can sustain before failure, relating applied stress to fatigue life through empirical laws.

2.1.1 Damage-accumulation processes

Table 2 summarises four widely used stochastic processes for modelling cumulative damage. In each case $X(t)$ denotes the cumulative damage at time t with $X(0) = 0$; all processes have independent, stationary increments (Lévy processes).

Table 2: Stochastic damage-accumulation models. Parameters: $\alpha, \beta > 0$ (Gamma shape rate / inverse scale), μ (drift), σ^2 (diffusion coefficient), λ (IG shape rate), ν (Poisson intensity), F_Y (jump-size distribution).

Model	Increment $X(t) - X(s)$	Marginal $X(t)$	Mon.	Closure under +
Gamma process [1]	$\text{Ga}(\alpha(t-s), \beta)$	$\text{Ga}(\alpha t, \beta)$	Yes	Same rate β
Wiener process [32]	$\mathcal{N}(\mu(t-s), \sigma^2(t-s))$	$\mathcal{N}(\mu t, \sigma^2 t)$	No	Always
Inverse Gaussian process [34]	$\text{IG}(\mu(t-s), \lambda(t-s)^2)$	$\text{IG}(\mu t, \lambda t^2)$	Yes	Same μ^2/λ
Compound Poisson [13, 31]	$\sum_{n=1}^{N(t-s)} Y_n$	$\text{CP}(\nu t, F_Y)$	If $Y \geq 0$	Always

Gamma process. Introduced by Abdel-Hameed [1] and surveyed in depth by van Noortwijk [30], the Gamma process is the most widely used model for monotonic degradation (wear, corrosion, crack growth). The process $\{X(t)\}$ has non-negative independent increments $X(t) - X(s) \sim \text{Ga}(\alpha(t-s), \beta)$, where $\alpha > 0$ is the shape rate and $\beta > 0$ is the inverse scale (rate). Its mean and variance at time t are $\mathbb{E}[X(t)] = \alpha t / \beta$ and $\text{Var}[X(t)] = \alpha t / \beta^2$. The Gamma family is closed under convolution when the rate parameter β is shared: if $X_1 \sim \text{Ga}(\alpha_1, \beta)$ and $X_2 \sim \text{Ga}(\alpha_2, \beta)$ are independent, then $X_1 + X_2 \sim \text{Ga}(\alpha_1 + \alpha_2, \beta)$.

Wiener process (Brownian motion with drift). The Wiener degradation model [32] sets $X(t) = \mu t + \sigma W(t)$, where $W(t)$ is a standard Brownian motion, $\mu > 0$ the drift rate, and $\sigma > 0$ the diffusion coefficient. Increments are normally distributed: $X(t) - X(s) \sim \mathcal{N}(\mu(t-s), \sigma^2(t-s))$. The normal family is unconditionally closed under addition. Unlike the other three models, the Wiener process is *not* monotone: sample paths can decrease, making it suitable for degradation phenomena with short-term fluctuations. A notable property is that the first-passage time to a fixed threshold τ follows an Inverse Gaussian distribution: $T_\tau \sim \text{IG}(\tau/\mu, \tau^2/\sigma^2)$.

Inverse Gaussian process. Proposed as a degradation model by Ye and Chen [34], the IG process has non-negative increments $X(t) - X(s) \sim \text{IG}(\mu(t-s), \lambda(t-s)^2)$, with mean μt and variance $\mu^3 t / \lambda$. Its increments are heavier-tailed than those of the Gamma process, making it appropriate when occasional large damage jumps are observed. The IG distribution is closed under convolution when the dispersion parameter μ^2/λ is shared across summands; however, the sum of two independent IG *processes* with different parameters is not itself an IG process [34].

Compound Poisson process. The compound Poisson process $X(t) = \sum_{n=1}^{N(t)} Y_n$ models shock-based degradation [13], where shocks arrive according to a Poisson process $N(t)$ with rate ν and each shock causes a random damage Y_n drawn i.i.d. from a distribution F_Y . When $Y_n \geq 0$, the process is monotone. The sum of independent compound Poisson processes is always a compound Poisson process. This model is natural for discrete damage events (e.g. impacts, overloads) rather than continuous wear; see Wang [31] for a recent application to structural reliability.

2.1.2 Discrete-time formulation

The continuous-time processes of Table 2 are sampled at discrete steps in the fleet management setting of Section 1: each day k , the damage of component (i, ℓ) grows by a random increment $\Delta D_{i\ell k}^{(j)}$ whose distribution depends on the mission j served that day. The cumulative damage after k steps is therefore a random walk

$$D_{i\ell k} = D_{i\ell,0} + \sum_{t=1}^k \Delta D_{i\ell t}, \quad (14)$$

where the increments $\Delta D_{i\ell t}$ are independent but *not* identically distributed (different missions may be served on different days). Table 3 restates the four models in this discrete-time setting.

Table 3: Discrete-time specialisation of the damage-accumulation models. Each row gives the distribution of a single-step increment ΔD and the resulting marginal of $D_k = \sum_{t=1}^k \Delta D_t$ when all increments share the parameters required for closure.

Model	Increment ΔD_t	Marginal D_k (homogeneous)	Closure condition
Gamma	$\text{Ga}(\alpha_t, \beta)$	$\text{Ga}(\sum_t \alpha_t, \beta)$	Same rate β
Gaussian	$\mathcal{N}(\mu_t, \sigma_t^2)$	$\mathcal{N}(\sum_t \mu_t, \sum_t \sigma_t^2)$	Always
Inv. Gaussian	$\text{IG}(\mu_t, \lambda_t)$	$\text{IG}(\sum_t \mu_t, \frac{(\sum_t \mu_t)^2}{\sum_t (\mu_t^2 / \lambda_t)})$	Same μ^2 / λ
Compound Poisson	$\sum_{n=1}^{N_t} Y_n$	$\text{CP}(\sum_t \nu_t, F_Y)$	Same F_Y

The closure conditions determine whether the marginal distribution of $D_{i\ell k}$ remains within the same parametric family after summing increments from different missions. When the condition holds, the distribution of $D_{i\ell k}$ is available in closed form as a function of the cumulative parameters, which is essential for evaluating the reliability constraint $\Pr(D_{i\ell k} > \tau) \leq \varepsilon$ efficiently. When it does not hold—for instance, summing Gamma increments with different rate parameters—the marginal must be computed numerically or approximated.

In the fleet management problem of Section 1, each day brings an increment with mission-dependent parameters. Closure therefore requires that the parameter shared across missions is indeed constant for a given component: the rate β for Gamma, or the dispersion μ^2 / λ for Inverse Gaussian. The shape (or mean/variance) parameters may vary freely across missions; they simply accumulate additively. For the Gaussian model, no restriction is needed: normal random variables are unconditionally closed under addition.

2.1.3 Remaining-life models: the S–N curve and Palmgren–Miner rule

The stress–life (S–N) approach, originating with Wöhler’s experiments in the 1850s, relates the stress amplitude σ_a to the number of cycles to failure N_f . The most common analytical form is Basquin’s law [3]:

$$\sigma_a^m N_f = C, \quad (15)$$

where $m > 0$ is the fatigue exponent and $C > 0$ is a material constant, both typically fitted from experimental data. Equivalently, $N_f = C \sigma_a^{-m}$, which appears as a straight line on a log–log plot.

Under variable-amplitude loading, the Palmgren–Miner linear damage rule [20, 24] defines the cumulative damage fraction as

$$D_{\text{PM}} = \sum_i \frac{n_i}{N_{f,i}}, \quad (16)$$

where n_i is the number of applied cycles at stress level $\sigma_{a,i}$ and $N_{f,i}$ is the corresponding life from the S–N curve. Failure is predicted when $D_{\text{PM}} \geq 1$. The rule assumes that damage accumulates linearly and independently of load ordering—an approximation that has been questioned experimentally (the critical damage sum D_{PM}^* at failure is scattered around 1 rather than being exactly 1) and extended stochastically by treating D_{PM} and/or D_{PM}^* as random variables [6, 9, 15, 22, 29].

Two sources of randomness enter the Palmgren–Miner accumulation:

- (i) *S–N scatter*. At a given stress level σ_a , the fatigue life N_f is not a deterministic constant but a random variable whose distribution reflects material variability, surface finish, and testing conditions. The increment $\Delta D_{\text{PM}} = n/N_f$ therefore inherits this randomness.
- (ii) *Random critical damage sum*. Experimental evidence shows that the critical value D_{PM}^* at failure is itself scattered (typically in the range 0.5–2.0) rather than being exactly 1 [29]. This can be modelled by replacing the deterministic threshold with a random one, or equivalently by normalising the damage by a random capacity.

In the fleet management framework of Section 1, source (i) is the relevant one: the threshold is fixed at $\tau = 1$ and the randomness enters through the increments. We therefore focus on the distributional assumptions for N_f and the resulting distribution of ΔD_{PM} . Table 4 summarises the three most common choices.

Table 4: Distributional assumptions for S–N scatter and their effect on the Palmgren–Miner increment $\Delta D_{\text{PM}} = n/N_f$. In all cases, $\mu(\sigma_a)$ and the scale/shape parameters depend on the stress level through Basquin’s law (15).

Model	Life N_f	Increment ΔD_{PM}	Mon.	Closure under +
Log-normal	$\ln N_f \sim \mathcal{N}(\mu, \sigma^2)$	$\text{LogN}(\ln n - \mu, \sigma^2)$	Yes	No
Weibull	$N_f \sim \text{Wei}(k, \lambda)$	Inv. Weibull (Fréchet)	Yes	No
Gumbel	$\ln N_f \sim \text{Gumbel}(\mu, \beta)$	Log-Gumbel	Yes	No

Log-normal. The most widely used assumption [22]: at a fixed stress level σ_a , the fatigue life satisfies $\ln N_f \sim \mathcal{N}(\mu(\sigma_a), \sigma^2)$, where $\mu(\sigma_a) = \ln C - m \ln \sigma_a$ from Basquin’s law and σ captures the scatter. Since $\ln(n/N_f) = \ln n - \ln N_f$, each increment is log-normally distributed: $\Delta D_{\text{PM}} \sim \text{LogN}(\ln n - \mu, \sigma^2)$, with mean $n e^{-\mu+\sigma^2/2}$ and variance $n^2 e^{-2\mu+\sigma^2} (e^{\sigma^2} - 1)$. The log-normal family is *not* closed under summation: the sum of independent log-normals has no standard parametric form. In practice, $D_{\text{PM},k} = \sum_{t=1}^k \Delta D_{\text{PM},t}$ is often approximated as log-normal via moment-matching (Fenton–Wilkinson method), but this is an approximation whose quality degrades as the number of summands grows or the coefficient of variation increases.

Weibull. Widely adopted in fatigue standards (e.g. Eurocode, ASTM): $N_f \sim \text{Wei}(k, \lambda(\sigma_a))$ with shape k and scale λ depending on the stress level. The reciprocal $1/N_f$ follows an inverse Weibull (Fréchet) distribution, so $\Delta D_{\text{PM}} = n/N_f$ is a scaled Fréchet variable. The sum of independent Fréchet variables has no closed form, making the marginal of $D_{\text{PM},k}$ intractable analytically.

Gumbel (extreme-value type I). Castillo and Fernández-Canteli [9] proposed a compatible statistical model for the entire S–N field based on a Gumbel (minimum) distribution for $\ln N_f$: $\ln N_f \sim \text{Gumbel}(\mu(\sigma_a), \beta)$. The increment $\Delta D_{\text{PM}} = n/N_f$ then follows a log-Gumbel distribution. As with the other two families, the log-Gumbel is not closed under summation.

Contrast with damage-accumulation models. A fundamental difference with the processes of Table 2 is that *none* of the standard distributional assumptions for S–N scatter yields increments that are closed under summation. The marginal distribution of $D_{\text{PM},k}$ after k loading blocks must therefore always be computed via convolution (analytically intractable in general), Monte Carlo simulation, or moment-based approximations. In the fleet management formulation, this means that the reliability constraint $\Pr(D_{\text{PM},k} \geq 1) \leq \varepsilon$ cannot be evaluated in closed form from the cumulative parameters alone, in contrast to the Gamma or Gaussian models where the CDF of the marginal is directly available. The upper bounds of Section 2.2 (Markov, Cantelli, Chernoff) remain applicable: because the mean and variance of $D_{\text{PM},k}$ are additive over increments regardless of the distributional family, the Cantelli bound provides a conservative but tractable constraint. Alternatively, if the scatter is small enough that a deterministic S–N curve (e.g. the mean or a design percentile) is an acceptable approximation, the problem reduces to the deterministic Miner’s rule and all constraints become linear.

2.1.4 Conservative surrogates for the empirical damage-increment distribution

In the remaining-life setting the per-mission increment is the Palmgren–Miner contribution $\Delta D^{(j)} = \sum_i n_i/N_{f,i}$ of Equation (16), and in our intended use its distribution is not postulated but *measured*: a rainflow count [22] of load histories collected for mission j yields an empirical probability density of the non-negative increment $\Delta D^{(j)}$. We deliberately do *not* fit a parametric law to this density—a fit ties the model to one data set and one functional form, undermining the generality the abstract formulation seeks—and instead summarise it by a few distributional functionals (moments, support, selected quantiles) and replace it by a *conservative* surrogate. Conservative here means that the surrogate never *under*-states damage, so that the reliability constraint $\Pr(D_{i\ell k} > \tau) \leq \varepsilon$ and the loop constraint (10) remain valid by construction. Two features of the problem delimit the admissible options. First, both constraints act on the *accumulated* damage $D = \sum_t \Delta D_t$ of Equation (14), not on a single increment, so the surrogate must be evaluated after summation. Second, as observed in the contrast above, none of the standard S–N scatter families is closed under summation, so the law of D is generally unavailable in closed form and must be either bounded or replaced by a convolution-closed surrogate.

Crucially, Route I never requires the moments or support of the accumulated damage D in advance: it consumes only *per-mission* information, which the schedule then assembles through the assignment counts. Fix a component (i, ℓ) and an accumulation window, write $D = D_{i\ell}$ for its damage over that window, and let $n_{ij} = \sum_t x_{ijt}$ be the number of days mission j is served. Modelling repeated executions of mission j on component ℓ as independent draws from a mission-specific increment law—with mean $\mu_{\ell j}$, variance $v_{\ell j}$, support $[0, b_{\ell j}]$, and cumulant generating function $\psi_{\ell j}(s) = \ln \mathbb{E}[e^{s\Delta D^{(j)}}]$, where $\mathbb{E}[e^{sX}]$ is the moment-generating function—independence across days gives

$$\mathbb{E}[D] = \sum_j n_{ij} \mu_{\ell j}, \quad \text{Var}[D] = \sum_j n_{ij} v_{\ell j}, \quad \ln M_D(s) = \sum_j n_{ij} \psi_{\ell j}(s). \quad (17)$$

Each functional of D is therefore a count-weighted sum of per-mission quantities—linear in the schedule, with the support entering only through the per-mission $b_{\ell j}$, never through a shared worst-case value. Repetition *merges* into the tail description additively: serving mission j a total of n_{ij} times multiplies its cumulant generating function by n_{ij} , and distinct missions combine by summation, so a gentle mission keeps its small $\psi_{\ell j}$ beside a punishing one rather than being averaged with it. These per-mission functionals support two routes. *Route I* bounds the upper tail of D directly from the functionals in (17). *Route II* replaces each increment by a stochastically larger surrogate drawn from a family closed under convolution and propagates the order to the sum, using the following standard closure property of the usual stochastic order $\leq_{\text{st}}$ (Definition 2.2).

Proposition 2.1 (Closure of $\leq_{\text{st}}$ under independent convolution). *Let $\Delta D_t \leq_{\text{st}} \widetilde{\Delta D}_t$ for $t = 1, \dots, k$, with $\{\Delta D_t\}_t$ mutually independent and $\{\widetilde{\Delta D}_t\}_t$ mutually independent. Then $\sum_t \Delta D_t \leq_{\text{st}} \sum_t \widetilde{\Delta D}_t$. More generally, the conclusion holds with the sum replaced by any coordinatewise non-decreasing map g (see Shaked and Shanthikumar [28, §1.A]).*

Proof. By the coupling characterisation of $\leq_{\text{st}}$ (Definition 2.2), for each t there exist $\hat{X}_t \stackrel{d}{=} \Delta D_t$ and $\hat{Y}_t \stackrel{d}{=} \widetilde{\Delta D}_t$ on a common probability space with $\hat{X}_t \leq \hat{Y}_t$ almost surely. Realise these couplings independently across t . Then $\sum_t \hat{X}_t \leq \sum_t \hat{Y}_t$ almost surely, while independence gives $\sum_t \hat{X}_t \stackrel{d}{=} \sum_t \Delta D_t$ and $\sum_t \hat{Y}_t \stackrel{d}{=} \sum_t \widetilde{\Delta D}_t$; hence $\sum_t \Delta D_t \leq_{\text{st}} \sum_t \widetilde{\Delta D}_t$. Any non-decreasing g preserves the almost-sure inequality, yielding the general statement. $\square$

Proposition 2.1 is what makes Route II usable beyond a single sum: the dynamics (9) is a coordinatewise non-decreasing map of the increments and the (independent) reset variables, so a per-increment stochastic upper bound propagates to a bound on the entire state D_{ilk} —repairs and replacements included—and hence to $\Pr(D_{ilk} > \tau)$ at every step. Table 5 collects the resulting options together with their guarantees and trade-offs.

The Route I entries are elementary, which is why they require no distributional assumption. With $D \geq 0$ and $\tau > 0$, Markov’s inequality $\Pr(D > \tau) \leq \mathbb{E}[D]/\tau$ follows from $\tau \mathbf{1}\{D > \tau\} \leq D$. Applying Markov to $(D - \mathbb{E}[D] + c)^2$ and optimising over $c > 0$ gives the one-sided Chebyshev (Cantelli) bound, for $\tau > \mathbb{E}[D]$,

$$\Pr(D > \tau) \leq \frac{\text{Var}[D]}{\text{Var}[D] + (\tau - \mathbb{E}[D])^2}. \quad (18)$$

Both are the *tight* bounds attainable from the first one and two moments, respectively [5], and both are attained (for the closed event $D \geq \tau$) by two-point distributions—which is exactly why they ignore, and are loose under, the independent-sum structure. The exponential (Cramér–Chernoff) bound,

$$\Pr(D > \tau) \leq \inf_{s>0} e^{-s\tau} M_D(s) = \inf_{s>0} \exp\left(-s\tau + \sum_t \ln M_{\Delta D_t}(s)\right), \quad (19)$$

follows from Markov applied to e^{sD} and is valid for every fixed $s > 0$ [10]; Bernstein’s inequality is the specialisation for independent summands bounded in $[0, b]$, $\Pr(D - \mathbb{E}[D] > a) \leq \exp\left(-\frac{a^2/2}{\text{Var}[D] + ba/3}\right)$ with $a = \tau - \mathbb{E}[D]$ [7]. Unlike Markov and Cantelli, these use the additive log-MGF of (17) and therefore exploit independence.

Per-mission inputs do not flatten the missions. Collapsing the summands to a single worst-case support $b = \max_j b_{\ell j}$ would make Bernstein treat every mission as the most damaging one, erasing the differences the schedule exploits; this flattening is an artefact of the closed form, not of Route I. Kept per term, each bound retains every mission’s own support and tail. Hoeffding reads $\Pr(D - \mathbb{E}[D] > a) \leq \exp(-2a^2/\sum_j n_{ij}b_{\ell j}^2)$ with $a = \tau - \mathbb{E}[D]$, where mission j enters

Table 5: Ways to summarise the empirical per-mission damage-increment distribution and propagate it to the accumulated damage $D = \sum_t \Delta D_t$ for the constraint $\Pr(D > \tau) \leq \varepsilon$. “Conservative” means a guaranteed upper bound on the failure probability (Route I) or $\leq_{\text{st}}$ dominance (Route II); “ $\tilde{D}$ ” means the method yields an explicit surrogate *law* for the sum—needed, e.g., for the loop constraint, which compares two survival functions—as opposed to a one-sided bound. All inputs are *per mission*; the schedule assembles them through the assignment counts n_{ij} via Equation (17), so the support enters only per mission and is never shared.

Method	Inputs	Cons.?	$\tilde{D}$?	Pros	Cons	Ref.
<i>Route I — bound the tail of D directly (no surrogate law)</i>						
Markov	mean	Yes	No	needs only the mean; always valid	extremely loose	[5]
Cantelli	mean, var	Yes	No	only two moments; tight as a two-moment bound	ignores independence; loose in the tail	[5]
Bernstein	per-mission mean, var, support b_j	Yes	No	exploits independence; few parameters; moment inputs linear in the counts	needs per-mission support; one-sided only	[7]
Chernoff	per-mission c.g.f. $\psi_j(s)$ (or mean + support)	Yes	No	tightest in the tail; count-weighted; valid for any fixed $s > 0$	needs the c.g.f.; empirical c.g.f. is an estimate, not a bound	[10]
<i>Route II — dominate each increment, then convolve (explicit $\tilde{D}$)</i>						
Moment-matched Gamma / IG	mean, var	No	Yes	explicit law; closed under +; ordering tables apply	matches moments but does not dominate (survival curves cross)	[30, 34]
Dominating shared-rate Gamma	moments + empirical survival	Yes	Yes	conservative <i>and</i> explicit; serves reliability and loop; shape linear in the schedule	shared rate + dominance margin add slack; commits to Gamma	[28]

only through its own $b_{\ell j}^2$, and the Chernoff bound (19) becomes the count-weighted reliability constraint

$$\Pr(D > \tau) \leq \inf_{s>0} \exp\left(-s\tau + \sum_j n_{ij} \psi_{\ell j}(s)\right), \quad (20)$$

in which each mission contributes its own cumulant generating function $\psi_{\ell j}$. For a fixed s the exponent is linear in the assignment counts, so (20) is a linear reliability constraint; the inner infimum couples s to the counts—for a fixed schedule it is a one-dimensional convex problem whose minimiser s^* solves $\sum_j n_{ij} \psi'_{\ell j}(s) = \tau$ (the exponentially tilted mean of D equals τ), but since s^* depends on the counts one either fixes a conservative s (any $s > 0$ remains a valid bound) or sweeps a small grid.

Bernstein's inequality is the one Route I bound that resists this per-mission treatment, and the reason is the same coupling seen from the other side: its closed form requires a *single* support $b_\ell = \max_j b_{\ell j}$ in the correction term (only the variance stays per-mission). Explicitly, with $\kappa = \ln(1/\varepsilon)$, the reliability constraint $\Pr(D > \tau) \leq \varepsilon$ becomes

$$\sum_j n_{ij} \mu_{\ell j} + \frac{\kappa b_\ell}{3} + \sqrt{\left(\frac{\kappa b_\ell}{3}\right)^2 + 2\kappa \sum_j n_{ij} v_{\ell j}} \leq \tau, \quad (21)$$

a constraint whose left-hand side is *concave* in the assignment counts (equivalently $(\kappa b_\ell/3)^2 + 2\kappa \sum_j n_{ij} v_{\ell j} \leq (\tau - \kappa b_\ell/3 - \sum_j n_{ij} \mu_{\ell j})^2$ with the bracket non-negative): by (17) the mean and the variance are *affine* in the counts, so the square root is concave and the feasible set is reverse-convex rather than second-order-cone-representable—in the classical SOCP chance constraint the variance is a convex quadratic of the decisions, whereas here independence across days makes it affine, which flips the curvature. For the MILP this is a feature rather than an obstacle: every tangent line to the concave square root *over*-estimates it, so replacing the square root in (21) by a fixed tangent—or by the piecewise-linear pointwise minimum of several tangents, encoded exactly as a PWL equality (SOS2)—yields linear constraints in the counts that imply (21) and therefore remain conservative. Retaining a per-mission support $b_{\ell j}$ in the correction is possible only by keeping the per-term cumulant sum $\sum_j n_{ij} (s^2 v_{\ell j}/2)/(1 - s b_{\ell j}/3)$ and minimising over s —which is exactly the Chernoff/Bennett bound (20). The s -free closed form (21) and per-mission support resolution therefore cannot be obtained simultaneously.

Guaranteed bound versus estimate. For (20) to be a true bound, the per-mission ingredient must be an *upper* bound on $\psi_{\ell j}$. Two per-mission choices suffice and keep the support local: the mean-and-support bound $\psi_{\ell j}(s) \leq \ln(1 + (\mu_{\ell j}/b_{\ell j})(e^{s b_{\ell j}} - 1))$, attained by the two-point law on $\{0, b_{\ell j}\}$; or the empirical cumulant generating function $\hat{\psi}_{\ell j}(s) = \ln(\frac{1}{m} \sum_{r=1}^m e^{s \Delta D_r^{(j)}})$ built from the m recorded executions of mission j , which is an *estimate*, not a bound, and so must carry a one-sided confidence margin. Repeating a mission tightens both ingredients—more executions sharpen $\hat{\psi}_{\ell j}$, $\mu_{\ell j}$ and $b_{\ell j}$ and shrink the margin—which is the precise sense in which undergoing a mission many times refines its tail description.

The choice among the Route I bounds depends on how deep into the tail the constraint sits. In the regime relevant to a reliability constraint with small ε —several standard deviations above $\mathbb{E}[D]$ —the exponential bounds are markedly tighter: in a representative numerical experiment (non-negative increments bounded in $[0, 1]$, $k \in \{30, 100\}$ summands), at $\varepsilon \sim 10^{-4}$ the Chernoff bound stayed within roughly one order of magnitude of the empirical tail, whereas the Cantelli bound over-stated it by more than two orders; Bernstein fell between them and overtook Cantelli only as the tail deepened and the number of summands grew.¹ Cantelli, though the simplest, is therefore best reserved for short accumulation windows or as a coarse always-valid fallback.

¹Recorded as a numerical illustration, not a proof; the test code and output are kept in the verification workspace. Asymptotically the gap is structural: Cantelli decays only polynomially, as $\text{Var}[D]/(\tau - \mathbb{E}[D])^2$, whereas (19) decays exponentially in the deviation.

Route II yields not a bound but an explicit surrogate law. Matching only the first two moments of a convolution-closed family (Gamma, with shape $\alpha = \mathbb{E}[\Delta D]^2 / \text{Var}[\Delta D]$ and rate $\beta = \mathbb{E}[\Delta D] / \text{Var}[\Delta D]$; or inverse Gaussian) is convenient—the sum stays in the family under the closure conditions of Table 3—but is *not* conservative: a moment-matched law generally crosses the empirical survival function rather than dominating it. Conservatism is recovered by using more of the empirical information than its moments. Fix a common rate β^* and, per mission, inflate the shape α_j until the surrogate survival dominates the empirical one, $\overline{\text{Ga}}(x; \alpha_j, \beta^*) \geq \widehat{F}_j(x)$ for all x —always achievable, since $\overline{\text{Ga}}(x; \alpha, \beta^*) \rightarrow 1$ as $\alpha \rightarrow \infty$. By Proposition 2.1 the accumulated surrogate is then $\widetilde{D} = \text{Ga}(A, \beta^*)$ with cumulative shape $A = \sum_{t,j} x_{ijt} \alpha_j$ linear in the schedule, the reliability constraint reduces to the scalar condition $A \leq A^*(\varepsilon)$, and the loop constraint to $A_{2H} \leq A_H$ exactly as in Equation (26). The price is the slack introduced by the shared rate and the dominance margin.

In summary, for the one-sided reliability constraint we use the count-weighted Chernoff bound (20)—which keeps each mission’s own tail—or its per-term Hoeffding specialisation when only per-mission means and supports are trusted; when an explicit surrogate law is required—notably for the loop constraint, a comparison of two survival functions—we use the dominating shared-rate Gamma of the last row, which lands on the tractable Gamma machinery of Section 2.2. The irreducible trade-off is that the bounds assume nothing about the shape but yield only a tail bound, whereas the dominating Gamma yields a propagatable law at the cost of a family commitment and some conservatism.

2.1.5 The loop constraint under the Route-I bounds

The loop constraint (10) requires $\Pr(D_{2H} > \tau) \leq \Pr(D_H > \tau)$ so that concatenating the schedule keeps the failure probability below ε indefinitely. Under Route I we never observe these probabilities; we have only an upper bound $U_k \geq \Pr(D_k > \tau)$ from the count-weighted Chernoff bound (20) or the Bernstein constraint (21). The tempting translation—replace both sides by their bounds and impose $U_{2H} \leq U_H$ —is *not* sufficient for repeatability, and it is worth seeing precisely why.

Why $U_{2H} \leq U_H$ does not chain. The bound U is a many-to-one function of the state’s parameters, and its level sets are not invariant under the dynamics. Bernstein’s bound, for instance, depends on $(\mathbb{E}[D], \text{Var}[D])$ through a single scalar, so $U_{2H} = U_H$ can hold with $(\mathbb{E}[D_{2H}], \text{Var}[D_{2H}]) \neq (\mathbb{E}[D_H], \text{Var}[D_H])$ —a larger mean offset by a smaller variance, say. One further period of the dynamics maps these two distinct parameter points to *different* bound values, so $U_{3H} > U_{2H} = U_H$ is possible: the certified bound can creep upward over repetitions, and the single-period inequality does not induct.² Adding a fixed margin, $U_{2H} \leq U_H - c$, does not repair this, because the defect is the non-invariance of the level sets, not a constant gap.

The certificate stays valid at every period. This is not an undetectable worsening. The functionals Route I actually tracks—the moments $\mathbb{E}[D_k], \text{Var}[D_k]$, or the cumulant generating function—are exact functionals of the true state and are *closed* under the dynamics (9): increments add them, imperfect repair scales them, replacement resets them. Consequently U_k is a valid bound at *every* period, $\Pr(D_{kH} > \tau) \leq U_{kH}$, and the true probability can exceed ε only if U_{kH} itself does. Repeatability therefore needs a loop condition that forces $U_{kH} \leq \varepsilon$ for all k , which means imposing it on a statistic that both determines U and propagates monotonically—not on the lossy scalar U .

²A concrete instance: $\tau = 10$, $b_\ell = 3$, $\rho = 0.5$, repeated block = one repair day followed by mission days with $\sum_j n_j \mu_{\ell j} = 3$ and $\sum_j n_j v_{\ell j} = 3.99$; from $(\mathbb{E}[D_H], \text{Var}[D_H]) = (4, 9)$ the Cantelli bound gives $U_H = 0.2000$ and $U_{2H} = 0.1997 \leq U_H$, yet $U_{3H} = 0.2151 > U_H$, with $U_{mH} \rightarrow 0.2495$. An analogous Bernstein instance (block variance 3.15) violates at $m = 3$ as well. The moment condition (22) rejects both, since $\mathbb{E}[D_{2H}] = 5 > 4$. Test scripts are kept in the verification workspace.

Looping on the moments (the native condition). Impose the loop on the exactly tracked moments,

$$\mathbb{E}[D_{2H}] \leq \mathbb{E}[D_H] \quad \text{and} \quad \text{Var}[D_{2H}] \leq \text{Var}[D_H]. \quad (22)$$

Under the dynamics, the mean and variance evolve by an increment ($\mathbb{E} += \sum_j n_{ij} \mu_{\ell_j}$, $\text{Var} += \sum_j n_{ij} v_{\ell_j}$), by imperfect repair ($\mathbb{E} \mapsto (1 - \rho) \mathbb{E}$, $\text{Var} \mapsto (1 - \rho)^2 \text{Var}$), or by replacement (reset to the constants of D^{new}); each operation is non-decreasing in $\mathbb{E}$ and in Var separately, so the two moments evolve as independent monotone recursions. Componentwise domination therefore chains: writing F for the once-per-period (second-half) map, (22) states $F(\cdot)$ does not increase either moment, whence $(\mathbb{E}[D_{(m+1)H}], \text{Var}[D_{(m+1)H}]) \leq (\mathbb{E}[D_{mH}], \text{Var}[D_{mH}])$ for every $m \geq 1$. Because the Bernstein bound (21)—and likewise the Cantelli bound (18)—is increasing in both $\mathbb{E}[D]$ and $\text{Var}[D]$, the certified risk is then non-increasing across periods, $U_{kH} \leq U_H \leq \varepsilon$, and the schedule is repeatable. Both inequalities in (22) are linear in the tracked moment states $(\mathbb{E}[D_k], \text{Var}[D_k])$, which the model carries as auxiliary variables (the moment recursion couples the maintenance binaries to these states through products, linearised exactly by standard big- M constructions), so the loop adds two linear constraints per component; they are strictly stronger than the single relation $U_{2H} \leq U_H$, and that strengthening is exactly what makes the induction valid.

Looping on the cumulant, and why imperfect repair forbids it. If reliability is instead certified by a fixed- s Chernoff bound, the loop may be placed on the cumulant,

$$\sum_j n_{ij}^{(2H)} \psi_{\ell_j}(s_0) \leq \sum_j n_{ij}^{(H)} \psi_{\ell_j}(s_0), \quad (23)$$

where $n_{ij}^{(k)}$ is the post-reset cumulative count of mission j up to day k . This chains *provided the repeated block contains only additive increments and full replacement*: there $M_D(s_0)$ propagates multiplicatively, $M_{D+\Delta D}(s_0) = M_D(s_0) M_{\Delta D}(s_0)$, so it is a self-propagating scalar and (23) induct. Imperfect (ARD) repair breaks this. With $D^+ = (1 - \rho)D^-$,

$$M_{D^+}(s_0) = \mathbb{E}[e^{s_0(1-\rho)D^-}] = M_{D^-}((1 - \rho)s_0) : \quad (24)$$

the repaired state's transform at s_0 is the old transform at the *shifted* argument $(1 - \rho)s_0$, which is not determined by $M_{D^-}(s_0)$ —the repair slides the state along the transform curve. The single scalar $M(s_0)$ is then no longer self-contained, the count-weighted cumulant ceases to be additive across the repair, and (23) no longer chains. The moments escape this because scaling acts on each of them by a fixed power, $\mathbb{E} \mapsto (1 - \rho)\mathbb{E}$ and $\text{Var} \mapsto (1 - \rho)^2 \text{Var}$, keeping their recursion closed—which is why (22) is the form to use whenever imperfect repair appears inside the repeated block.

Summary. Under Route I the loop constraint must be imposed on a sufficient, self-propagating statistic, never on the scalar bound U . The componentwise moment condition (22) is the native choice: two linear constraints, valid for the Bernstein and Cantelli certificates and for every maintenance operator in (9), imperfect repair included. The cumulant condition (23) is available when the certificate is the fixed- s Chernoff bound and the repeated block is free of imperfect repair. In either case the guarantee is non-worsening of the *certified* risk, hence $\Pr(D_{kH} > \tau) \leq \varepsilon$ for all concatenations—the operative meaning of repeatability for a bound-based model.

2.2 Stochastic orderings

Comparing the degradation state of a component at different points in the schedule requires a notion of “one random variable being larger than another.” Several stochastic orders formalise this idea with increasing strength [28, 33]. Let X and Y be two non-negative random variables

with CDFs F_X, F_Y , survival functions $\bar{F}_X = 1 - F_X, \bar{F}_Y = 1 - F_Y$, and densities f_X, f_Y (when they exist).

Definition 2.2 (Usual stochastic order). $X \leq_{\text{st}} Y$ if $F_X(t) \geq F_Y(t)$ for all $t \in \mathbb{R}$, or equivalently $\Pr(X > t) \leq \Pr(Y > t)$ for all t .

Definition 2.3 (Hazard rate order). $X \leq_{\text{hr}} Y$ if $\bar{F}_Y(t)/\bar{F}_X(t)$ is non-decreasing in t over $\{t : \bar{F}_X(t) > 0\}$. Equivalently, the hazard rates satisfy $r_X(t) \geq r_Y(t)$ for all t , where $r(t) = f(t)/\bar{F}(t)$.

Definition 2.4 (Likelihood ratio order). $X \leq_{\text{lr}} Y$ if $f_Y(t)/f_X(t)$ is non-decreasing in t over the union of their supports.

These three orders are progressively stronger [28, Ch. 1]. The implications are strict—none of the converses holds in general:

$$X \leq_{\text{lr}} Y \implies X \leq_{\text{hr}} Y \implies X \leq_{\text{st}} Y \implies \mathbb{E}[X] \leq \mathbb{E}[Y]. \quad (25)$$

Remark 2.5 (Connection to the loop constraint). The loop constraint (10) requires $\Pr(D_{il,2H} > \tau) \leq \Pr(D_{il,H} > \tau)$, which is a *pointwise* comparison of the survival functions at the single threshold τ . This is strictly weaker than $D_{il,2H} \leq_{\text{st}} D_{il,H}$, which demands the inequality at every point. Establishing the usual stochastic order is therefore a convenient sufficient condition for the loop constraint, though not a necessary one.

2.2.1 Stochastic ordering within each degradation family

For the fleet management problem, the relevant comparison is between two damage states of the same component—for instance the state at mid-horizon versus end-of-horizon—both belonging to the same parametric family. Tables 6–8 give, for each family from Table 2, the necessary and sufficient conditions for each of the three orderings. An entry marked “—” indicates that no simple closed-form characterisation is known to the author; in those cases the $\leq_{\text{lr}}$ conditions remain a valid sufficient condition via (25).

Table 6: Conditions for $X \leq_{\text{lr}} Y$ (likelihood ratio order) within each family. All conditions are necessary and sufficient.

Family	X	Y	Condition
Gamma	$\text{Ga}(\alpha_1, \beta_1)$	$\text{Ga}(\alpha_2, \beta_2)$	$\alpha_1 \leq \alpha_2$ and $\beta_1 \geq \beta_2$
Gaussian	$\mathcal{N}(\mu_1, \sigma_1^2)$	$\mathcal{N}(\mu_2, \sigma_2^2)$	$\mu_1 \leq \mu_2$ and $\sigma_1 = \sigma_2$
Inv. Gaussian	$\text{IG}(\mu_1, \lambda_1)$	$\text{IG}(\mu_2, \lambda_2)$	$\lambda_2 \geq \lambda_1$ and $\lambda_1/\mu_1^2 \geq \lambda_2/\mu_2^2$
Comp. Poisson	$\text{CP}(\nu_1, F_{Y_1})$	$\text{CP}(\nu_2, F_{Y_2})$	$\nu_1 \leq \nu_2, F_{Y_1} = F_{Y_2}, Y \geq 0$

Table 7: Conditions for $X \leq_{\text{hr}} Y$ (hazard rate order) within each family.

Family	X	Y	Condition
Gamma	$\text{Ga}(\alpha_1, \beta_1)$	$\text{Ga}(\alpha_2, \beta_2)$	$\alpha_1 \leq \alpha_2$ and $\beta_1 \geq \beta_2$
Gaussian	$\mathcal{N}(\mu_1, \sigma_1^2)$	$\mathcal{N}(\mu_2, \sigma_2^2)$	$\mu_1 \leq \mu_2$ and $\sigma_1 = \sigma_2$
Inv. Gaussian	$\text{IG}(\mu_1, \lambda_1)$	$\text{IG}(\mu_2, \lambda_2)$	— (suf.: $\leq_{\text{lr}}$ conditions)
Comp. Poisson	$\text{CP}(\nu_1, F_{Y_1})$	$\text{CP}(\nu_2, F_{Y_2})$	— (suf.: $\leq_{\text{lr}}$ conditions)

Table 8: Conditions for $X \leq_{\text{st}} Y$ (usual stochastic order) within each family.

Family	X	Y	Condition
Gamma	$\text{Ga}(\alpha_1, \beta_1)$	$\text{Ga}(\alpha_2, \beta_2)$	$\alpha_1 \leq \alpha_2$ and $\beta_1 \geq \beta_2$
Gaussian	$\mathcal{N}(\mu_1, \sigma_1^2)$	$\mathcal{N}(\mu_2, \sigma_2^2)$	$\mu_1 \leq \mu_2$ and $\sigma_1 = \sigma_2$
Inv. Gaussian	$\text{IG}(\mu_1, \lambda_1)$	$\text{IG}(\mu_2, \lambda_2)$	— (suf.: $\leq_{\text{lr}}$ conditions)
Comp. Poisson	$\text{CP}(\nu_1, F_{Y_1})$	$\text{CP}(\nu_2, F_{Y_2})$	— (suf.: $\leq_{\text{lr}}$ conditions)

Gamma. The log-likelihood ratio of $\text{Ga}(\alpha_2, \beta_2)$ over $\text{Ga}(\alpha_1, \beta_1)$ has derivative $(\alpha_2 - \alpha_1)/x - (\beta_2 - \beta_1)$, which is non-negative for all $x > 0$ if and only if $\alpha_2 \geq \alpha_1$ and $\beta_2 \leq \beta_1$. Conversely, if $\alpha_1 > \alpha_2$ then $F_X(t) \sim (\beta_1 t)^{\alpha_1} / (\alpha_1 \Gamma(\alpha_1)) \ll (\beta_2 t)^{\alpha_2} / (\alpha_2 \Gamma(\alpha_2)) \sim F_Y(t)$ as $t \rightarrow 0^+$, violating $\leq_{\text{st}}$; if $\beta_1 < \beta_2$ then $\bar{F}_X(t)/\bar{F}_Y(t) \rightarrow \infty$ as $t \rightarrow \infty$, again violating $\leq_{\text{st}}$. Hence all three orders coincide for the Gamma family [28, Sec. 1.C]. Under the closure condition (same β), they all reduce to $\alpha_1 \leq \alpha_2$.

Gaussian. The CDF of $\mathcal{N}(\mu, \sigma^2)$ is $\Phi((t - \mu)/\sigma)$. The inequality $\Phi((t - \mu_1)/\sigma_1) \geq \Phi((t - \mu_2)/\sigma_2)$ for all t requires $(t - \mu_1)/\sigma_1 \geq (t - \mu_2)/\sigma_2$ for all t , which is linear in t and can hold identically only if $\sigma_1 = \sigma_2$, in which case it reduces to $\mu_1 \leq \mu_2$. The likelihood ratio is then $\exp((\mu_2 - \mu_1)x/\sigma^2 + \text{const})$, non-decreasing in x . Thus all three orders coincide: $\sigma_1 = \sigma_2$ and $\mu_1 \leq \mu_2$.

Inverse Gaussian. The log-likelihood ratio of $\text{IG}(\mu_2, \lambda_2)$ over $\text{IG}(\mu_1, \lambda_1)$ has derivative

$$\frac{\lambda_1}{2\mu_1^2} - \frac{\lambda_2}{2\mu_2^2} + \frac{\lambda_2 - \lambda_1}{2x^2}.$$

This is non-negative for all $x > 0$ if and only if $\lambda_2 \geq \lambda_1$ (from $x \rightarrow 0^+$) and $\lambda_1/\mu_1^2 \geq \lambda_2/\mu_2^2$ (from $x \rightarrow \infty$). Under the closure condition (same μ^2/λ), these simplify to $\mu_1 \leq \mu_2$. For the general two-parameter case, closed-form necessary and sufficient conditions for $\leq_{\text{hr}}$ and $\leq_{\text{st}}$ do not appear in the standard references; the $\leq_{\text{lr}}$ conditions serve as a sufficient criterion.

Compound Poisson. With the same jump-size distribution F_Y ($Y \geq 0$), $\text{CP}(\nu_1, F_Y) \leq_{\text{lr}} \text{CP}(\nu_2, F_Y)$ when $\nu_1 \leq \nu_2$: a higher arrival rate produces stochastically more shocks and hence more cumulative damage. When the jump-size distributions differ or are not constrained to be non-negative, the ordering conditions depend on the specific form of F_Y and no universal closed-form characterisation is available.

Implications for the fleet problem. Under the closure conditions of Table 3, the cumulative damage $D_{i\ell k}$ depends on the schedule only through the sum of the mission-dependent shape (or mean) parameters. Stochastic ordering of two damage states therefore reduces to comparing scalar cumulative parameters: for Gamma, the total shape $\sum_t \alpha_t$; for Gaussian, the total mean $\sum_t \mu_t$ (the total variance $\sum_t \sigma_t^2$ is equal if the same missions are served in different order, or is irrelevant in the Wiener-process formulation where σ^2 is a fixed component parameter); for IG, the total mean $\sum_t \mu_t$. This scalar comparison is what makes the loop constraint (10) tractable within each model family.

2.2.2 Loop constraint for each degradation model

We now specialise the loop constraint (10) to each family. Let D_H and D_{2H} denote the degradation of a single component at the mid-point and end of the horizon, with parameters deter-

mined by the schedule and any intervening maintenance. The constraint reads $\Pr(D_{2H} > \tau) \leq \Pr(D_H > \tau)$. We drop the (i, ℓ) subscripts for clarity.

Gamma. Let $D_H \sim \text{Ga}(A_H, \beta)$ and $D_{2H} \sim \text{Ga}(A_{2H}, \beta)$, where A_k denotes the effective cumulative shape at step k (accounting for repairs and replacements). The survival function is $\Pr(D > \tau) = Q(A, \beta\tau)$, where $Q(a, x) = \Gamma(a, x)/\Gamma(a)$ is the regularised upper incomplete gamma function. Since $Q(a, x)$ is strictly increasing in a for fixed $x > 0$, the loop constraint becomes a scalar inequality:

$$A_{2H} \leq A_H. \quad (26)$$

This is a linear constraint on the cumulative shape parameters, directly expressible in terms of the scheduling and maintenance decision variables.

Gaussian. Let $D_H \sim \mathcal{N}(\mu_H, \sigma_H^2)$ and $D_{2H} \sim \mathcal{N}(\mu_{2H}, \sigma_{2H}^2)$. The survival function is $\Pr(D > \tau) = \Phi((\mu - \tau)/\sigma)$, where Φ is the standard normal CDF. Since Φ is strictly increasing, the loop constraint is equivalent to

$$\frac{\mu_{2H} - \tau}{\sigma_{2H}} \leq \frac{\mu_H - \tau}{\sigma_H}, \quad (27)$$

or equivalently $(\tau - \mu_{2H})/\sigma_{2H} \geq (\tau - \mu_H)/\sigma_H$. When $\sigma_H = \sigma_{2H}$ (e.g. if the Wiener process uses a fixed per-step diffusion coefficient and the same number of mission days is served in each half), this reduces to $\mu_{2H} \leq \mu_H$. In general, condition (27) couples the mean and standard deviation: a higher variance at $2H$ can be compensated by a sufficiently lower mean.

Inverse Gaussian. Let $D_H \sim \text{IG}(\mu_H, \lambda_H)$ and $D_{2H} \sim \text{IG}(\mu_{2H}, \lambda_{2H})$. The CDF of $\text{IG}(\mu, \lambda)$ admits the closed form

$$F_{\text{IG}}(x; \mu, \lambda) = \Phi\left(\sqrt{\frac{\lambda}{x}}\left(\frac{x}{\mu} - 1\right)\right) + e^{2\lambda/\mu} \Phi\left(-\sqrt{\frac{\lambda}{x}}\left(\frac{x}{\mu} + 1\right)\right). \quad (28)$$

The loop constraint becomes

$$F_{\text{IG}}(\tau; \mu_{2H}, \lambda_{2H}) \geq F_{\text{IG}}(\tau; \mu_H, \lambda_H). \quad (29)$$

This is an analytical but nonlinear inequality in the four parameters $(\mu_H, \lambda_H, \mu_{2H}, \lambda_{2H})$. Under the closure condition (same μ^2/λ), the $\leq_{\text{lr}}$ results of Table 6 guarantee that $\mu_{2H} \leq \mu_H$ is a sufficient condition.

Compound Poisson. When $D \sim \text{CP}(\nu, F_Y)$, the survival function does not admit a universal closed form. Conditioning on the number of shocks gives the series

$$\Pr(D > \tau) = \sum_{n=0}^{\infty} \frac{e^{-\nu} \nu^n}{n!} \Pr\left(\sum_{k=1}^n Y_k > \tau\right), \quad (30)$$

which can be evaluated numerically (it converges rapidly) or bounded.

Upper bounds when no closed form is available. When the survival function $\Pr(D > \tau)$ cannot be computed in closed form—as for the compound Poisson process, or when the closure condition fails and the marginal of the sum is not within a standard family—the loop constraint can still be enforced through conservative upper bounds on the failure probability. Three standard choices, in order of increasing tightness, are:

- (i) *Markov bound.* $\Pr(D > \tau) \leq \mathbb{E}[D]/\tau$. Requires only the mean; very loose but always valid.

(ii) *Cantelli bound* (one-sided Chebyshev). For $\tau > \mathbb{E}[D]$:

$$\Pr(D > \tau) \leq \frac{\text{Var}[D]}{\text{Var}[D] + (\tau - \mathbb{E}[D])^2}. \quad (31)$$

Requires the first two moments; tighter than Markov when τ is well above the mean.

(iii) *Chernoff bound*. $\Pr(D > \tau) \leq \inf_{s>0} e^{-s\tau} M_D(s)$, where $M_D(s) = \mathbb{E}[e^{sD}]$ is the moment-generating function. For the compound Poisson process, $M_D(s) = \exp(\nu(M_Y(s) - 1))$, which is available in closed form whenever M_Y is. This yields the tightest exponential bound but requires solving a one-dimensional optimisation over s .

In an optimisation setting, the Cantelli bound is often the most practical: the loop constraint reduces to a constraint on $\mathbb{E}[D_{2H}]$ and $\text{Var}[D_{2H}]$ relative to $\mathbb{E}[D_H]$ and $\text{Var}[D_H]$, both of which are additive over increments and therefore linear (or low-order polynomial) in the scheduling decisions.

2.3 Maintenance strategies

Maintenance actions on a repairable component range between two extremes: *minimal repair*, which restores the component to the state it was in just before failure (“as bad as old”, ABAO), and *perfect repair* or *replacement*, which resets it to a like-new condition (“as good as new”, AGAN). Real-world interventions almost always fall between these two extremes; they are collectively referred to as *imperfect maintenance* [25]. This section surveys the principal modelling families.

2.3.1 Taxonomy of repair actions

Table 9 classifies five canonical maintenance actions by their effect on the component state.

Table 9: Taxonomy of maintenance actions, from least to most restorative.

Action	Effect on damage D	Model
No intervention	$D^+ = D^-$	Damage unchanged
Minimal repair (ABAO)	$D^+ = D^-$	Failure intensity restored to pre-failure level; damage history preserved [8]
Imperfect repair	$D^- > D^+ > D^{\text{new}}$	Partial reduction of damage or virtual age (see below)
Perfect repair (AGAN)	$D^+ = D^{\text{new}}$	Component returned to like-new state
Replacement	$D^+ = D^{\text{new}}$	Physical component substituted; equivalent to perfect repair

Here D^- and D^+ denote the damage immediately before and after maintenance, and D^{new} is the damage of a new component.

2.3.2 Virtual age models

Kijima [18] formalised imperfect repair through the concept of *virtual age*. After the n -th repair, the component behaves as if it had age V_n (rather than its calendar age), and the failure intensity at calendar time t is $\lambda(t - T_n + V_n)$, where T_n is the time of the n -th repair and $\lambda(\cdot)$ is the baseline intensity. Two variants are standard:

- **Kijima I.** The repair acts only on the damage accumulated since the previous repair:

$$V_n = V_{n-1} + (1 - \rho) (T_n - T_{n-1}),$$

where $\rho \in [0, 1]$ is the repair efficiency. Only the most recent inter-repair period is partially rejuvenated (“memory-one”).

- **Kijima II.** The repair acts on the entire accumulated virtual age:

$$V_n = (1 - \rho) (V_{n-1} + T_n - T_{n-1}).$$

The full history is rejuvenated (“infinite memory”).

In both cases $\rho = 0$ gives minimal repair (ABAO) and $\rho = 1$ gives perfect repair (AGAN).

2.3.3 Reduction-based models

Doyen and Gaudoin [12] proposed a unifying taxonomy by distinguishing the *quantity* that the repair reduces: the accumulated damage (or virtual age), or the failure intensity. Each class admits a memory-one variant (subscript 1) and an infinite-memory variant (subscript ∞).

Arithmetic Reduction of Damage / Age (ARD, ARA). The repair removes a fraction ρ of the accumulated damage:

$$\text{ARD}_1 : D^+ = D^- - \rho (D^- - D^{\text{last}}), \quad (32)$$

$$\text{ARA}_1 : V^+ = V^- - \rho (V^- - V^{\text{last}}), \quad (33)$$

$$\text{ARD}_\infty : D^+ = (1 - \rho) D^-, \quad (34)$$

$$\text{ARA}_\infty : V^+ = (1 - \rho) V^-, \quad (35)$$

where D^{last} (resp. V^{last}) is the damage (resp. virtual age) at the end of the previous maintenance epoch. ARD_∞ coincides with Kijima II expressed in damage units, and ARA_1 coincides with Kijima I expressed in virtual age.

Arithmetic Reduction of Intensity (ARI). Instead of reducing the virtual age, the repair directly reduces the failure intensity:

$$\text{ARI}_1 : \lambda^+ = \lambda^- - \rho (\lambda^- - \lambda^{\text{last}}), \quad (36)$$

$$\text{ARI}_\infty : \lambda^+ = (1 - \rho) \lambda^-. \quad (37)$$

ARI models are natural when maintenance targets the failure rate (e.g. condition-based interventions) rather than the underlying degradation state. They are in general not equivalent to ARA/ARD models: a reduction in intensity does not correspond to a unique reduction in virtual age unless the baseline intensity is strictly monotone [12].

2.3.4 Connection to the fleet management formulation

The abstract repair operator $\mathcal{R}(D^-, \rho)$ introduced in Section 1 encompasses ARD, and replacement models. In particular:

- ARD_∞ : $\mathcal{R}(D^-, \rho) = (1 - \rho) D^-$ (used in the scoping as the default “ARD” policy);
- ARD_1 : $\mathcal{R}(D^-, \rho) = D^- - \rho (D^- - D^{\text{last}})$, requiring the additional state variable D^{last} ;
- Replacement / AGAN: $\mathcal{R}(D^-, \rho) = D^{\text{new}}$ (captured by the separate replacement variable $x_{i\ell k}^r$ in the formulation).

The choice of repair model affects the degradation dynamics (9), the definition of the repair cost variable $z_{i\ell k}$, and the tractability of the resulting optimisation problem. ARI models are not directly compatible with the damage-tracking formulation of Section 1 and would require reformulation in terms of the failure intensity rather than the damage state.

Maintenance under the Palmgren–Miner damage model. The remaining-life framework of Section 2.1.2 represents damage as the normalised fraction $D_{\text{PM}} = \sum_i n_i/N_{f,i} \in [0, 1]$, with failure at $D_{\text{PM}} \geq 1$. Because D_{PM} is additive over loading blocks, it fits the same damage-tracking structure as the stochastic processes of Table 2: each mission day contributes an increment $\Delta D_{\text{PM}} = n_j/N_{f,j}$ (deterministic if the S–N parameters are fixed, or random if scatter is modelled [22]). The repair operator $\mathcal{R}$ and the replacement action therefore apply directly:

- *Replacement* resets the fatigue fraction to zero: $D_{\text{PM}}^+ = 0$ (equivalently $D^{\text{new}} = 0$). This corresponds to substituting the physical component with a new one whose cycle count is zero.
- ARD_∞ : $D_{\text{PM}}^+ = (1 - \rho) D_{\text{PM}}^-$. Physically, this models an intervention that partially restores fatigue life (e.g. surface treatment, shot peening, localised crack repair) without resetting the full loading history.
- ARD_1 : $D_{\text{PM}}^+ = D_{\text{PM}}^- - \rho(D_{\text{PM}}^- - D_{\text{PM}}^{\text{last}})$. Only the damage accumulated since the last maintenance epoch is partially reversed.

A notable difference from the stochastic process models is that D_{PM} is bounded above by construction (failure at 1), so the threshold is fixed at $\tau = 1$ and the reliability constraint becomes $\Pr(D_{\text{PM}} \geq 1) \leq \varepsilon$. In the deterministic variant of Miner’s rule this is a hard feasibility check; in the stochastic extension it becomes a chance constraint analogous to (13f).

Closure under summation is trivially satisfied: D_{PM} is a sum of real-valued increments, so the marginal at any step is the convolution of the individual increments regardless of their distributional family. In the deterministic variant of Miner’s rule, comparing two schedules reduces to comparing two real numbers and no stochastic ordering is needed. In the stochastic extension (scatter in $N_{f,i}$), the parametric ordering conditions of Section 2.2 apply only if the distributional assumption on the scatter places the increments $n_i/N_{f,i}$ in one of the families of Table 2; otherwise the general ordering definitions still hold but must be verified case by case.

Remaining items for this section.

- ✓ Fatigue models (either damage accumulation or remaining life)
- Degradation-aware fleet management
- ✓ Maintenance strategies
- Stochastic orderings
- Similar applications
- Novelties of the project and its placement in the literature

3 Focus on electric vehicles

A battery-electric road vehicle—a transit bus or a passenger car—is a complex repairable system whose parts degrade at very different rates and through very different physical mechanisms. Not every part deserves an explicit place in the fleet-management model of Section 1. We retain a component if it satisfies two criteria:

- (C1) **Usage-driven degradation.** Its condition deteriorates measurably as a function of *how* the vehicle is used, so that the mission assigned on a given day produces a quantifiable damage increment $\Delta D_{ilk}^{(j)}$ (Section 1). A part that ages only with calendar time, independently of the mission served, carries no scheduling information and need not enter the per-mission degradation state.

(C2) **Operational significance.** Its failure or preventive renewal is a material driver of at least one of *safety*, *availability*, or *life-cycle maintenance cost*—the three quantities the schedule trades off.

We deliberately do *not* require that the component admit a degradation model expressible within the families of Section 2.2: a component can be worth considering even when the abstract model of Section 1 does not yet capture it. Such cases are listed here together with the modelling extension they would demand, rather than being silently dropped.

The components meeting these criteria fall into two broad groups. The *electrochemical and electrical powertrain*—traction battery, electric motor, power electronics, charging interface—degrades through electrochemical ageing and thermal and electrical cycling, driven by energy throughput, charge and discharge rates, and temperature. The *mechanical running gear*—tires, friction brakes, wheel bearings, suspension, reduction gear—degrades through wear and fatigue, driven by distance travelled, vehicle mass, road quality, and the acceleration/braking profile. Table 10 lists the principal candidates in each group, the dominant degradation mechanism, the operational quantity that drives it, and the reason the component is operationally significant.

Two domain-specific facts shape the relative importance of these components. First, the traction battery dominates: a recent predictive-maintenance study of electric buses centres entirely on the battery pack as the primary maintenance-critical component, models its end-of-life through the accumulated charge throughput, and shows that optimising the driving profile alone can cut the modelled capacity loss by 25% [17, Secs. 2, 4.2]—direct evidence that battery damage is steerable by exactly the kind of mission-assignment decisions our framework optimises. Second, regenerative braking absorbs most of the service braking, with estimated reductions of roughly 60–90% in friction-brake wear emissions [14, Sec. 6.1]; the relative wear burden consequently shifts to the tires, whose wear is aggravated by the higher mass of electric vehicles—to the point that, on motorways, the added tire-wear particle emissions are not offset by the regenerative-braking savings [14, Sec. 6.1].

Power-electronic and electronic components. The traction chain contains four electrical sub-families beyond the battery, and they fit the framework to differing degrees. The *power semiconductors* of the inverter and the DC–DC and on-board-charger *converters* are governed by junction-temperature-swing fatigue and map cleanly onto the remaining-life family (treated in detail below). The *converter capacitors*—DC-link and filtering—wear out by electrolyte vaporisation and dielectric ageing driven by ripple-current self-heating; this is a monotone, usage-driven accumulation (Arrhenius in temperature, power law in ripple current) and therefore also fits the framework, with the mission acting through the delivered power and the resulting ripple and thermal load.³ The *control and power printed circuit boards*—gate-driver, battery-management, and converter boards—are the borderline case: their solder-joint and plated-through-via fatigue is genuinely usage-driven (thermal cycling, power-on hours, vibration) and so satisfies (C1)–(C2), but a substantial part of their failures are sudden and effectively random rather than a smooth accumulation crossing a threshold. The damage-accumulation dynamics of Section 1 do not fully capture this mixed wear-out-plus-random behaviour; representing it faithfully would require either a compound-Poisson shock term (Section 2.2) or a failure-intensity formulation rather than the damage-state formulation of Equation (9). We list PCBs here, with this caveat, precisely so the limitation is recorded rather than hidden by the choice of model.

Modelling shortlist. For a first concrete instantiation of the abstract framework, the most informative components are those that are simultaneously safety- or cost-critical *and* equipped with mature, monotone degradation models that fit the closure and ordering analysis of Section 2.2: the *traction battery pack*, the *tires*, the *brake friction elements*, the *wheel bearings*, the

³The capacitor and PCB degradation-model forms in this paragraph are stated from standard power-electronics reliability theory; located primary sources have not yet been secured for them (see `inaccessible_papers.md`).

Table 10: Candidate electric-vehicle components for degradation-aware modelling. Relevance flags: **S** safety-critical, **A** availability driver (its failure or renewal takes the vehicle out of service), **C** maintenance-cost driver. The “mission driver” column names the operational quantity that determines the per-mission damage increment $\Delta D_{ilk}^{(j)}$.

Component	Dominant degradation mechanism	Mission driver	Rel.
<i>Electrochemical / electrical powertrain</i>			
Traction battery pack	Capacity fade and impedance growth: SEI growth, lithium plating, loss of active material (cycle + calendar ageing)	Energy throughput, DOD, C-rate, temperature	A, C
Traction motor	Winding-insulation thermal ageing; bearing wear	Thermal/load cycles, current	A, C
Power semiconductors (IGBT / SiC modules in the inverter and converters)	Thermal-cycling fatigue (bond-wire lift-off, solder cracking)	Power throughput, thermal cycles	A, C
Converter capacitors (DC-link, filtering)	Electrolyte vaporisation / dielectric wear-out (ripple-current self-heating)	Ripple current, temperature, operating hours	A, C
Control / power PCBs (gate drivers, BMS, converter boards)	Solder-joint and via thermo-mechanical fatigue; random infant/wear-out failures	Thermal cycles, power-on hours, vibration	A
Charging interface (inlet, connector)	Contact wear and contact-resistance growth over mating cycles	Number of charging events	A
<i>Mechanical: running gear and chassis</i>			
Tires	Tread abrasion (friction work at the contact patch)	Distance, vehicle mass, driving style, road surface	S, C
Brake friction elements (pads, discs)	Abrasive and thermal wear (mitigated by regenerative braking)	Number and energy of friction-braking events	S, C
Wheel bearings	Rolling-contact fatigue spalling	Distance, load	S, A
Suspension (springs, dampers, bushings)	Fatigue; loss of damping	Load cycles, road quality	A
Reduction gear / drive shafts	Wear, contact fatigue, lubricant ageing	Transmitted torque, distance	A
Auxiliaries (HVAC, battery thermal management, doors)	Wear, cycle-driven failures	Operating hours, actuation cycles	A

traction-motor insulation, and the *inverter and converter power modules*. Each maps onto a single per-component degradation variable D_{ilk} whose mission-dependent increment $\Delta D_{ilk}^{(j)}$ is governed by the operational driver in the third column of Table 10. A long highway mission, for instance, produces large tire-wear and deep-discharge battery increments, whereas a stop-intensive urban mission multiplies braking events and charge–discharge micro-cycles; this mission-dependence is exactly the coupling that the scheduling problem of Section 1 exploits.

Worked example: traction-battery capacity fade. The battery pack is the costliest component of an electric vehicle and the focus of its predictive maintenance [17, Secs. 2, 4.2], so it serves as the reference instantiation of the abstract increment $\Delta D_{ilk}^{(j)}$. Its degradation state is the relative capacity fade (one minus state-of-health), and the failure threshold τ is the maximum admissible fade (the end-of-life criterion). Mechanistically, capacity is lost through three side-reaction families—SEI film growth, lithium plating, and gas generation—whose contributions split into loss of lithium inventory and loss of active material [19, Secs. 2.1, 4.1, Eqs. 28–29]. At the fleet level this is abstracted by semi-empirical models that decompose the fade into a *calendar* term and a *cycling* term: the calendar term grows with an Arrhenius factor $a \exp(-E_a/RT)$ in temperature and a power law in time and state-of-charge, while the cycling term grows with the number of cycles modulated by temperature, C-rate, and depth-of-discharge stress factors [16, Eqs. 1–3]. An even simpler engineering abstraction links end-of-life directly to the accumulated charge throughput [17, Sec. 4.2]. For our framework the decisive structure is: the cycling fade is monotone and additive over days, and the daily increment is a function of the energy served on the mission—including the recharge that restores it, so that fast charging (high C-rate) makes the same mission costlier in damage—together with its depth-of-discharge and temperature profile. This is precisely a mission-dependent accumulation process of the Gamma or inverse-Gaussian type of Section 2.2; the non-stationary (power-law) time dependence is absorbed by the time-varying increment parameters that the formulation already allows ($\Delta D_{ilk}^{(j)}$ may depend on k). The calendar term, by contrast, accrues even on idle and maintenance days and is therefore *not* captured by the dynamics of Section 1, which leave the state unchanged on such days; over short horizons it can be neglected, while over long horizons the dynamics (9) must be extended with a small activity-independent drift added to every case, including idle. Maintenance also takes a specific form: replacing a subset of battery modules removes a corresponding fraction of the pack-level damage—a physical realisation of the ARD repair operator $\mathcal{R}$ with ρ the fraction of modules renewed—while full pack replacement is the AGAN action x^r .

3.1 Degradation models for the shortlisted components

The remaining shortlisted components degrade through the same two structural mechanisms identified for the battery—monotone usage-driven *wear/accumulation* and *cyclic-load fatigue life consumption*—and therefore map onto the same two model families of Section 2.2. Table 11 summarises the governing model for each, the operational quantity that drives the per-mission increment, and the abstract family it instantiates.

Tires. Tread loss is governed by the frictional work expended at the tire–road contact patch: the classical theory of Schallamach and Turner assumes abrasion proportional to the frictionally dissipated energy [26, Eq. 7] and derives that, at small slip, wear grows approximately with the *square* of the slip [26, Eq. 19]—so the wear rate rises steeply with the tangential forces generated by traction, braking, and cornering. The principal factors are the tire characteristics (compound, construction), the road-surface characteristics, and the vehicle operation [14, Sec. 2.1.2]. The higher kerb mass of electric vehicles raises these contact forces—enough that, on motorways, the added tire-wear particle emissions are not offset by the brake-wear savings of regenerative braking [14, Sec. 6.1]. Tread depth decreases monotonically from its new value to the legal

Table 11: Degradation models for the shortlisted electric-vehicle components and their mapping onto the abstract model families of Section 2.2. “Accumulation” denotes a monotone damage-accumulation process (Gamma / inverse-Gaussian / compound-Poisson); “remaining-life” denotes a cycles-to-failure model accumulated by the Palmgren–Miner rule.

Component	Degradation model / law	Mission driver	Abstract family
Battery pack	Calendar + cycle capacity fade: Arrhenius in T , power law in t/SOC , cycle-count term [16, Eqs. 1–3]; mechanistic SEI / plating / LAM [19, Sec. 2.1]	Energy throughput, DOD, C-rate, temperature	Accumulation (k -dependent)
Tires	Friction-work abrasion [26, Eq. 7]; factors: tire build, road surface, vehicle operation [14, Sec. 2.1.2]	Distance, mass, accel./braking, road surface	Accumulation
Brake friction pad	Archard-type wear per braking event [27, Eqs. 14, 22]; regenerative-braking mitigation [14, Sec. 6.1]	Number and energy of friction-braking events	Accumulation
Wheel bearing	Subsurface-fatigue spall propagation, continuum damage [23, Sec. 2.1.3]	Distance, speed, load	Accumulation / remaining-life
Traction-motor insulation	Arrhenius–Dakin thermal-endurance life [11]; Montsinger rule [21]; cumulative-damage extension [35, Eqs. 1–2]	Thermal / load cycles, current	Remaining-life (Miner)
Inverter power module	Coffin–Manson / Norris–Landzberg cycles-to-failure in ΔT_j [2, Eqs. 3–5]; Bayerer/CIPS08 [4, Eq. 1]	Power throughput, thermal cycles	Remaining-life (Miner)

limit, which plays the role of the threshold τ ; the per-mission increment depends on distance, load, road surface, and driving aggressiveness. For bus fleets, retreading acts as a near-AGAN replacement of the wear state, while tire rotation merely redistributes wear across positions and is left outside the per-component model.

Brake friction elements. Friction-pad material loss is governed by the energy dissipated at the pad-disc interface during braking; a physics-based route applies Archard’s law on a fractal, elastoplastic contact to obtain a per-cycle wear depth and hence a brake-cycle life $N = H_0/(S_0 h_{\text{wear}})$, with H_0 the admissible wear depth and h_{wear} the wear per braking [27, Eqs. 14, 22]. In an electric vehicle the regenerative brake absorbs most service braking, cutting friction-brake wear by an estimated 60–90% [14, Sec. 6.1]; the residual increments concentrate in the events where regeneration is unavailable or insufficient—battery near full charge, very low temperature, emergency stops. The cumulative pad wear remains a sum of non-negative, mission-dependent increments—one per friction-braking event—i.e. a monotone accumulation process whose rate now also depends on the energy-management strategy.

Wheel bearings. Bearing degradation proceeds in two stages: a long rolling-contact-fatigue *initiation* phase followed by rapid *spall propagation* once a surface defect forms. Ohana et al. [23, Sec. 2.1.3] model the propagation phase with a continuum-damage-mechanics formulation in which a scalar damage variable accumulates per over-rolling cycle, $dD/dN = (\sigma_{\text{eff}}/[\sigma_r(1-D)])^m$, with growth driven by the tensile residual stresses generated at the spall edge by repeated rolling-element impacts; the model is intended as the basis for a physics-based prognostic tool that estimates remaining useful life once a critical spall size is reached. Because the underlying driver is the number of load cycles—and hence running distance and speed—the bearing damage is a monotone accumulation process, well represented by a Gamma or inverse-Gaussian process (initiation) feeding a remaining-life threshold (propagation).

Traction-motor insulation. The winding insulation is the principal life-limiting element of the traction motor, and its dominant ageing mechanism is thermal: elevated winding temperature accelerates the chemical degradation of the dielectric. Dakin recognised thermal ageing as a chemical reaction whose rate constant follows the Arrhenius exponential law $K = A e^{-B/T}$ in absolute temperature [11, Eq. 5], giving the thermal-endurance life $L(T) = a \exp(b/T)$ [35, Eq. 1]; the earlier empirical rule of Montsinger states that the ageing rate roughly doubles for each 8 °C temperature increase [21, Synopsis, p. 776]—a halving of life per fixed rise that Dakin showed to be the constant-coefficient approximation of the Arrhenius law, accurate at one temperature only [11, Fig. 1 and discussion of Eq. 5]. Under a mission-varying temperature profile, the fraction of life consumed in an interval dt at temperature $T_i(t)$ is $dt/L[T_i(t)]$, and the total consumed fraction is the integral of this loss rate over the profile—a cumulative-damage (Palmgren–Miner) extension of the Arrhenius model [35, Eq. 2 and Fig. 10]. The abstract description therefore coincides with the remaining-life model of Section 2.1.2, with the mission driving the thermal load (current, cooling, ambient); notably, the same review also surveys Wiener- and Gamma-process formulations of insulation degradation [35, Sec. 5], directly matching the stochastic families of Section 2.2.

Inverter power modules. The IGBT or SiC modules of the traction inverter fail predominantly by thermo-mechanical fatigue—aluminium bond-wire lift-off and solder-layer cracking—driven by the coefficient-of-thermal-expansion mismatch under repeated junction-temperature swings ΔT_j [4, Sec. 1]. The number of cycles to failure is classically captured by a Coffin–Manson power law $N_f = \alpha (\Delta T_j)^{-n}$ [2, Eq. 3], refined by adding an Arrhenius factor in the mean junction temperature T_{jm} , $N_f = A (\Delta T_j)^\alpha \exp(E_a/(k_b T_{jm}))$ [2, Eq. 4], and generalised by the multi-factor Bayerer (CIPS08) model $N_f = K \Delta T_j^{\beta_1} e^{\beta_2/(T_j+273)} t_{\text{on}}^{\beta_3} I^{\beta_4} V^{\beta_5} D^{\beta_6}$, fitted by

multiple regression over power-cycling tests spanning temperature swing, power-on time t_{on} , current per bond wire I , blocking voltage V , and bond-wire diameter D [4, Eq. 1]. Being purely empirical, the model is valid only within the fitted test range and is not to be extrapolated [4, Sec. 3.1]; its coefficients also require recalibration per device family—a recent re-fit on dedicated power-cycling tests shifted $(\beta_1, \beta_2, \beta_3)$ substantially and cut the predicted B10 inverter lifetime from 43 to 13 years [2, Abstract and Sec. 5]. Under a realistic mission mix the thermal cycles have heterogeneous amplitudes, and the accumulated damage is summed by the Palmgren–Miner rule $Q = \sum_i n_i/N_{f,i}$, failure at $Q = 1$ —placing the power module squarely in the remaining-life family of Section 2.1.2, with the mission determining the transmitted power and hence the temperature-swing spectrum.

Common structure. Every shortlisted component reduces to one of the two abstract descriptions of Section 2.2: a *monotone accumulation process* (battery cycling fade, tires, brakes, bearings) or a *remaining-life model accumulated by the Palmgren–Miner rule* (insulation, power modules). In both, the per-mission damage increment $\Delta D_{ilk}^{(j)}$ is set by an operational driver that the schedule controls—energy throughput and charging rate, distance and load, braking profile, or thermal cycling—which is precisely the structure the abstract formulation of Section 1 assumes. This confirms that the electric-vehicle fleet is a faithful concrete instance of the general problem, and fixes which degradation family parameterises each component when the model is specialised in Section 4.

Excluded parts. Components whose deterioration is essentially independent of the mission served—body corrosion, interior fittings, static low-voltage electronics, the 12-volt auxiliary battery—satisfy criterion (C2) but not (C1). They are best handled by time-based overhaul outside the optimisation and are therefore omitted from the per-mission degradation state, although their fixed maintenance windows can be imposed as exogenous availability constraints. The calendar-ageing component of the traction battery is the one borderline case and is handled by the drift extension discussed in the worked example above.

Remaining items for this section.

- ☐ Relevant components (mechanical and electrical) to model an electric vehicle
- ☐ Typical degradation model for each component, and the operational assumptions under which it holds

4 Optimization problem for a fleet of electric vehicles

Table 12 lists the additional symbols introduced when specialising the abstract formulation of Section 1 to a concrete fleet of electric vehicles. All symbols from Table 1 remain in force; the entries below either instantiate abstract quantities or add new ones.

Table 12: Additional notation for the electric-vehicle fleet specialisation (extends Table 1).

Symbol	Description
<i>Cost parameters</i>	
C_M	Cost of scheduling a maintenance day (scalar)
C_D	Damage regularisation coefficient (scalar)
$C_{R,il}$	Cost of imperfect repair per unit of expected degradation removed
$C_{\text{rep},il}$	Cost of full replacement of component ℓ of member i

4.1 Cost function

With the cost parameters of Table 12, the abstract objective of Section 1 takes the form:

$$\begin{aligned} \min \quad & C_M \sum_{k=1}^{2H} \sum_{i=1}^F x_{i0k} + \sum_{k=1}^{2H} \sum_{i=1}^F \sum_{\ell=1}^L C_{R,il} z_{ilk} + \sum_{k=1}^{2H} \sum_{i=1}^F \sum_{\ell=1}^L C_{\text{rep},il} x_{ilk}^r \\ & + C_D u. \end{aligned} \tag{38}$$

Each term is a weighted sum over the planning horizon: the first counts maintenance days, the second scales with the degradation removed by each repair, the third counts replacements, and the last penalises the worst-case aggregate expected damage via u (cf. (12)).

- Specialize the previous optimization problem to the case of electric vehicles
- Look for what can be exactly modeled and what no
- Provide the MILP approximation of such problem

5 Implementation details

- Provide all the project specifications for the proper implementation of the solving algorithm
- Provide the details to run an actual simulation of the system

References

- [1] Mohamed Abdel-Hameed. A gamma wear process. *IEEE Transactions on Reliability*, R-24(2):152–153, 1975. doi: 10.1109/TR.1975.5215123.
- [2] Omid Alavi, Ward De Ceuninck, and Michaël Daenen. Optimizing insulated-gate bipolar transistors’ lifetime estimation: A critical evaluation of lifetime model adjustments based on power cycling tests. *Energies*, 17(11):2616, 2024. doi: 10.3390/en17112616.
- [3] O. H. Basquin. The exponential law of endurance tests. *Proceedings of the American Society for Testing and Materials*, 10:625–630, 1910.
- [4] Reinhold Bayerer, Tobias Herrmann, Thomas Licht, Josef Lutz, and Marco Feller. Model for power cycling lifetime of IGBT modules—various factors influencing lifetime. In *5th International Conference on Integrated Power Electronics Systems (CIPS)*, Nuremberg, Germany, 2008.
- [5] Dimitris Bertsimas and Ioana Popescu. Optimal inequalities in probability theory: A convex optimization approach. *SIAM Journal on Optimization*, 15(3):780–804, 2005. doi: 10.1137/S1052623401399903.
- [6] Marta Böhm, Marek Kowalski, and Adam Nieslony. Non-linear probabilistic modification of Miner’s rule for damage accumulation. *Materials*, 14(23):7440, 2021. doi: 10.3390/ma14237440.
- [7] Stéphane Boucheron, Gábor Lugosi, and Pascal Massart. *Concentration Inequalities: A Nonasymptotic Theory of Independence*. Oxford University Press, 2013. ISBN 9780199535255.
- [8] Mark Brown and Frank Proschan. Imperfect repair. *Journal of Applied Probability*, 20(4): 851–859, 1983. doi: 10.2307/3213596.

- [9] Enrique Castillo and Alfonso Fernández-Canteli. *A Unified Statistical Methodology for Modeling Fatigue Damage*. Springer, 2009. doi: 10.1007/978-1-4020-9182-7.
- [10] Herman Chernoff. A measure of asymptotic efficiency for tests of a hypothesis based on the sum of observations. *The Annals of Mathematical Statistics*, 23(4):493–507, 1952. doi: 10.1214/aoms/1177729330.
- [11] Thomas W. Dakin. Electrical insulation deterioration treated as a chemical rate phenomenon. *Transactions of the American Institute of Electrical Engineers*, 67(1):113–122, 1948. doi: 10.1109/T-AIEE.1948.5059649.
- [12] Laurent Doyen and Olivier Gaudoin. Classes of imperfect repair models based on reduction of failure intensity or virtual age. *Reliability Engineering & System Safety*, 84(1):45–56, 2004. doi: 10.1016/S0951-8320(03)00173-X.
- [13] James D. Esary, Albert W. Marshall, and Frank Proschan. Shock models and wear processes. *The Annals of Probability*, 1(4):627–649, 1973. doi: 10.1214/aop/1176996891.
- [14] Julia C. Fussell, Meredith Franklin, David C. Green, Mats Gustafsson, Roy M. Harrison, William Hicks, Frank J. Kelly, Franceska Kishta, Mark R. Miller, Ian S. Mudway, Farzan Oroumiyeh, Liza Selley, Meng Wang, and Yifang Zhu. A review of road traffic-derived non-exhaust particles: Emissions, physicochemical characteristics, health risks, and mitigation measures. *Environmental Science & Technology*, 56(11):6813–6835, 2022. doi: 10.1021/acs.est.2c01072.
- [15] Hassan Hilow. Probabilistic Palmgren–Miner rule, with application to solder materials. *Journal of Materials Science: Materials in Electronics*, 28:12901–12913, 2017. doi: 10.1007/s10854-017-7117-8.
- [16] Md Sazzad Hosen, Joris Jaguemont, Joeri Van Mierlo, and Maitane Berecibar. Battery lifetime prediction and performance assessment of different modeling approaches. *iScience*, 24(2):102060, 2021. doi: 10.1016/j.isci.2021.102060.
- [17] Hisham Ibrahim, Ahmed M. Ali, and Tamer Attia. Optimizing electric bus performance via predictive maintenance: A combined experimental and modeling approach. *Frontiers in Future Transportation*, 5:1506866, 2025. doi: 10.3389/ffutr.2024.1506866.
- [18] Masaaki Kijima. Some results for repairable systems with general repair. *Journal of Applied Probability*, 26(1):89–102, 1989. doi: 10.2307/3214319.
- [19] Zhe Lv, Huinan Si, Zhe Yang, Jiawen Cui, Zhichao He, Lei Wang, Zhe Li, and Jianbo Zhang. Simplified mechanistic aging model for lithium ion batteries in large-scale applications. *Materials*, 18(6):1342, 2025. doi: 10.3390/ma18061342.
- [20] Milton A. Miner. Cumulative damage in fatigue. *Journal of Applied Mechanics*, 12(3):A159–A164, 1945. doi: 10.1115/1.4009458.
- [21] V. M. Montsinger. Loading transformers by temperature. *Transactions of the American Institute of Electrical Engineers*, 49(2):776–790, 1930. doi: 10.1109/T-AIEE.1930.5055572.
- [22] M. Muñiz-Calvente et al. A comparative review of time- and frequency-domain methods for fatigue damage assessment. *International Journal of Fatigue*, 2022.
- [23] Ravit Ohana, Renata Klein, Roni Shneck, and Jacob Bortman. New damage accumulation model for spall propagation mechanism in bearing raceways. *Materials*, 16(4):1750, 2023. doi: 10.3390/ma16041750.

- [24] Arvid Palmgren. Die Lebensdauer von Kugellagern. *Zeitschrift des Vereines Deutscher Ingenieure*, 68(14):339–341, 1924.
- [25] Hoang Pham and Hongzhou Wang. Imperfect maintenance. *European Journal of Operational Research*, 94(3):425–438, 1996. doi: 10.1016/S0377-2217(96)00099-9.
- [26] A. Schallamach and D. M. Turner. The wear of slipping wheels. *Wear*, 3(1):1–25, 1960. doi: 10.1016/0043-1648(60)90283-9.
- [27] Zhihua Sha, Qiang Hao, Jian Yin, Fujian Ma, Yu Liu, and Shengfang Zhang. Wear calculation and life prediction model of disc brake based on elastoplastic contact mechanics. *Advances in Mechanical Engineering*, 14(4), 2022. doi: 10.1177/16878132221093136.
- [28] Moshe Shaked and J. George Shanthikumar. *Stochastic Orders*. Springer Series in Statistics. Springer, 2007. doi: 10.1007/978-0-387-34675-5.
- [29] Chih-Kuang Shen. A statistical consideration of Miner’s rule. *International Journal of Fatigue*, 2(1):27–30, 1980. doi: 10.1016/0142-1123(80)90024-5.
- [30] Jan M. van Noortwijk. A survey of the application of gamma processes in maintenance. *Reliability Engineering & System Safety*, 94(1):2–21, 2009. doi: 10.1016/j.ress.2007.03.019.
- [31] Cao Wang. An explicit compound Poisson process-based shock deterioration model for reliability assessment of aging structures. *Journal of Traffic and Transportation Engineering (English Edition)*, 9(3):461–472, 2022. doi: 10.1016/j.jtte.2020.03.005.
- [32] G. A. Whitmore and Fred Schenkelberg. Modelling accelerated degradation data using Wiener diffusion with a time scale transformation. In *Lifetime Data Analysis*, volume 3, pages 27–45. 1997. doi: 10.1023/A:1009664101413.
- [33] Wikipedia contributors. Stochastic dominance — Wikipedia, the free encyclopedia, 2026. URL https://en.wikipedia.org/wiki/Stochastic_dominance. Accessed: 2026-05-21.
- [34] Zhi-Sheng Ye and Nan Chen. The inverse Gaussian process as a degradation model. *Technometrics*, 56(3):302–311, 2014. doi: 10.1080/00401706.2013.830074.
- [35] Jian Zhang, Jiajin Wang, Hongbo Li, Qin Zhang, Xiangning He, Cui Meng, Xiaoyan Huang, Youtong Fang, and Jianwei Wu. A review of reliability assessment and lifetime prediction methods for electrical machine insulation under thermal aging. *Energies*, 18(3):576, 2025. doi: 10.3390/en18030576.